

Natural Language Processing

Using tensorflow



BITTIGER

The Lifelong Learning Platform of Silicon Valley

Copyright Policy

All content included on the Site or third-party platforms as part of the class, such as text, graphics, logos, button icons, images, audio clips, video clips, live streams, digital downloads, data compilations, and software, is the property of BitTiger or its content suppliers and protected by copyright laws.

Any attempt to redistribute or resell will result in the appropriate legal action being taken.



We thank you in advance for respecting our copyrighted content.

Outline



1. Input Shape
2. RNN Feed
3. Model (variables and ops)

Input Shape



BITTIGER

The Lifelong Learning Platform of Silicon Valley

Shape



- First, concatenate all characters

千山鸟飞绝 ... 白华照寒水 ... 落叶满长安 ...



- `batch_len = data_len // batch_size` (`300 = 30010 // 100`)

Reshape



```
tf.reshape(raw_data[0:batch_size * batch_len],  
           [batch_size, batch_len])
```

Drop last 10 words

千山鸟飞绝 ... 白华照寒水 ... 落叶满长安

reshape

batch_size

batch_len

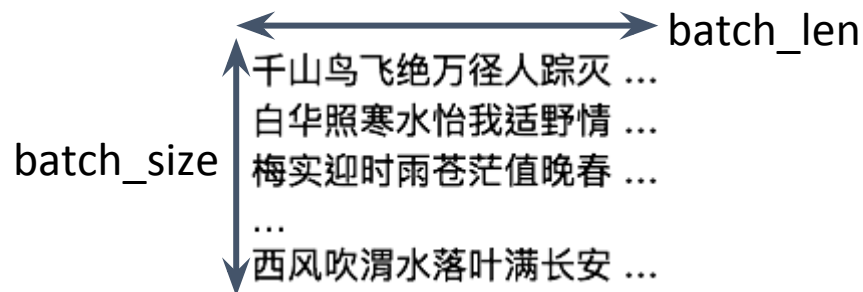
千山鸟飞绝万径人踪灭 ...
白华照寒水怡我适野情 ...
梅实迎时雨苍茫值晚春 ...
...
西风吹渭水落叶满长安 ...

Batch



- `epoch_size = (batch_len - 1) // num_steps`

$$59 = (300-1) // 5$$



```
i = tf.train.range_input_producer(epoch_size, shuffle=False).dequeue()  
i = 0, 1, 2, ..., 58
```

Batch



```
i = 0, 1, 2, ..., 58  
x = data[:, i*num_steps : (i+1)*num_steps]  
y = data[:, i*num_steps+1 : (i+1)*num_steps+1]
```

```
x1 = [0:100, 0*5 : 1*5]  
y1 = [0:100, 0*5+1 : 1*5+1]
```

千山鸟飞绝 万径人踪灭 ...
白华照寒水 怡我适野情 ...
梅实迎时雨 苍茫值晚春 ...
...
西风吹渭水 落叶满长安 ...

Batch



```
i = 0, 1, 2, ..., 58  
x = data[:, i*num_steps : (i+1)*num_steps]  
y = data[:, i*num_steps+1 : (i+1)*num_steps+1]
```

```
x2 = [0:100, 1*5 : 2*5]  
y2 = [0:100, 1*5+1 : 2*5+1]
```

千山鸟飞绝 万径人踪灭 ...
白华照寒水 怡我适野情 ...
梅实迎时雨 苍茫值晚春 ...
...
西风吹渭水 落叶满长安 ...

Batch



```
i = 0, 1, 2, ..., 58
x = data[:, i*num_steps : (i+1)*num_steps]
y = data[:, i*num_steps+1: (i+1)*num_steps+1]
```

```
                290:295
x59 = [0:100, 58*5:59*5]
y59 = [0:100, 58*5+1:59*5+1]
                291:296
```

千山鸟飞绝万径人踪灭 ...
白华照寒水怡我适野情 ...
梅实迎时雨苍茫值晚春 ...
...
西风吹渭水落叶满长安 ...

- Last part of data that cannot form a batch will be abandoned.



RNN Feed



BITTIGER

The Lifelong Learning Platform of Silicon Valley


Feed During Training



- Less information for the model when predicting > y2 words

- More information
- Simulate the condition where the model can predict the correct word

After Training



- When the model is well-trained, it should produce correct words on its own

Model: variables and ops



BITTIGER

The Lifelong Learning Platform of Silicon Valley

Variables and ops



- A layer may contain variables and operations

```
variable embedding_weights = tf.get_variable("embedding",  
                                             [vocab_size, embedding_size])  
  
op      embed_inputs =  
        tf.nn.embedding_lookup(embedding_weights, input)
```

- Or just op, no variable

```
tf.nn.dropout(embed_inputs, .....)
```

Can use the same name



- A layer may contain variables and operations

```
variable embedding_weights = tf.get_variable("embedding",  
                                             [vocab_size, embedding_size])  
  
op embed_inputs =  
    tf.nn.embedding_lookup
```

- Or just op, no variable


```
tf.nn.dropout(embed_inputs, ...)
```

You can use the same name to connect output of one layer to the input of another layer. (optional)

Step-by-step Building




- We will demonstrate with **variables** and **operations**
- Understand how to go from input to output



Building Blocks



1. Create `variable` if needed




```
embedding_weights =  
tf.get_variable("embedding",  
                [vocab_size, embedding_size])
```

Building Blocks



1. Create **variable** if needed
2. Add **op** that uses the variable




```
embedding_lookup(embedding_weights, input)
```

Building Blocks



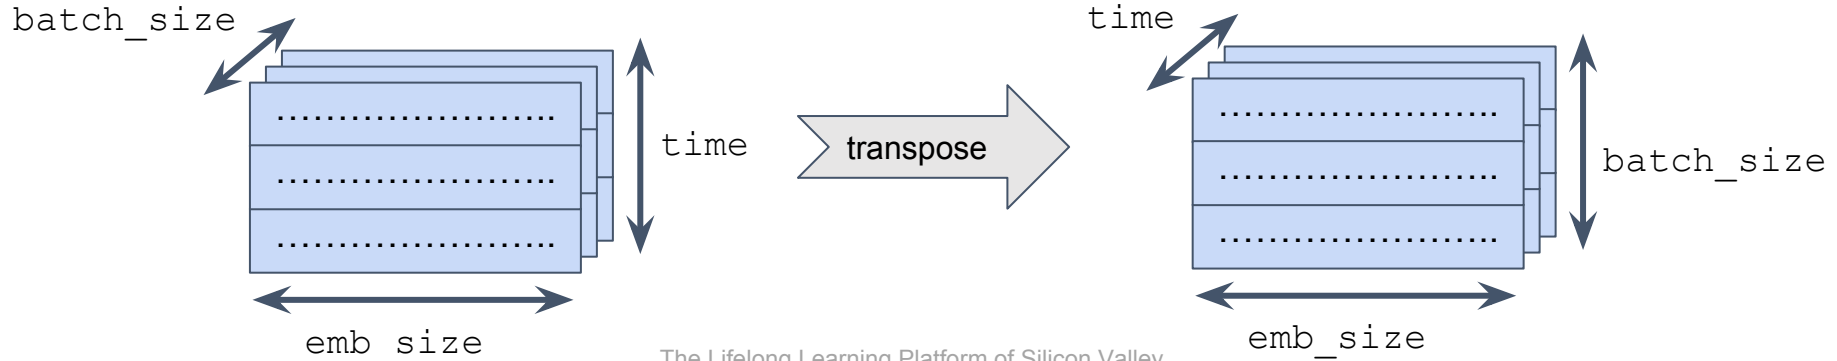
```
embedding_lookup(embedding_weights, [1, 30, 100])
```




Building Blocks




1. Create **variable** if needed
2. Add **op** that uses the variable
3. Transpose/reshape if needed before sending to the next layer



Building Blocks



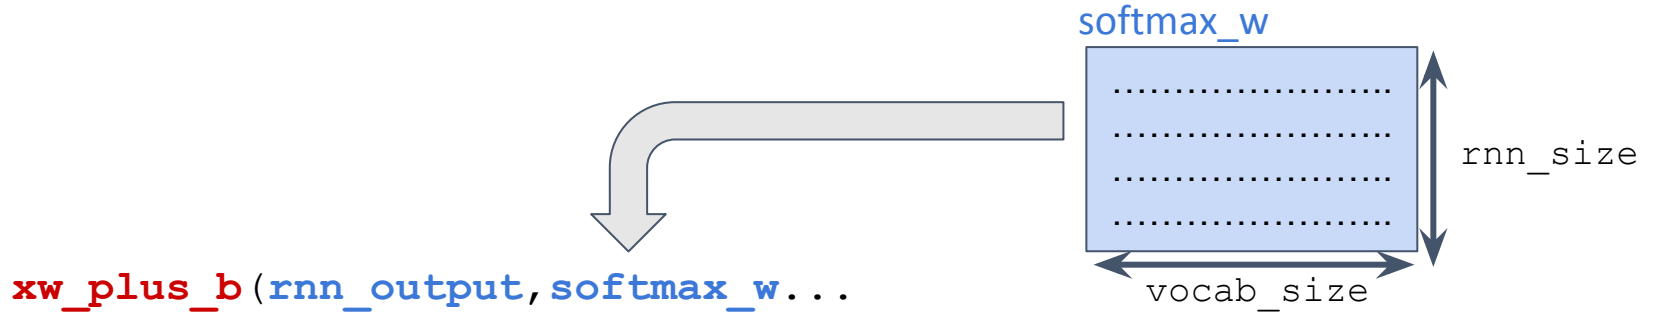
- Merge 1 and 2 dimension of RNN layer output




Building Blocks



1. Create `variable` if needed
2. Add `op` that uses the variable
3. Transpose/reshape if needed before sending to the next layer



Building Blocks



- Represent probability of word

A thick yellow square frame with a small gap on the right side.

Codelab



BITTIGER

The Lifelong Learning Platform of Silicon Valley

Attention



BITTIGER


The Lifelong Learning Platform of Silicon Valley

Concept




1. Calculate weight over encoder outputs
Using current decoder memory
2. Weighted sum of encoder outputs, “context vector”
3. Merge context with decoder output (word)
4. Send to decoder for next decoding step
5. Back to 1.


Attention




Attention




Attention




Attention



Attention



Attention



Definition



1. Calculate weight over encoder outputs

$$\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s) = \mathbf{v}_a^\top \tanh(\mathbf{W}_1 \mathbf{h}_t + \mathbf{W}_2 \bar{\mathbf{h}}_s)$$

2. Normalize weights

$$\alpha_{ts} = \frac{\exp(\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_s))}{\sum_{s'=1}^S \exp(\text{score}(\mathbf{h}_t, \bar{\mathbf{h}}_{s'}))}$$

3. Sum to get “context vector”

$$\mathbf{c}_t = \sum_s \alpha_{ts} \bar{\mathbf{h}}_s$$